



Hopscotch Hashing

18.11.2020

Sridhar M

2018111021

Monsoon 2020

IIIT Hyderabad

Overview

The project is on understanding and implementing Hopscotch hashing in the JAVA language.

Hopscotch hashing is a method of open addressing that builds on the prior work in chained hashing, linear probing, and Cuckoo hashing in order to design a new methodology. The key idea of hopscotch hashing is as follows:

- Assign a global neighborhood size H
- keep the hashed key within this neighborhood

Deliverables

1. Java implementation of Hopscotch hashing
2. Project report
3. Presentation of the project

Algorithmic overview

1. There is a single hash function h
2. Hashed item will always be found either in that home bucket or in one of the next $H-1$ neighboring bucket (hop)
3. Each bucket includes some hop-information using an H -bit bitmap (in code, represented by- **hope**), it indicates which of the items in the bucket's neighborhood hashed to the current bucket

Cache Locality

Hopscotch hashing has super-fast query times in practice is because of its excellent cache locality. Complexity for find operation is $O(1)$. Setting neighborhood size to be equal to the size of the cache line means there will be at most 2 cache misses before an object is found.

Comparison: Round Robin Hashing

- Hopscotch hashing is slightly more cache-friendly than Robin Hood variants but requires more complicated operations and implementations.

- Hopscotch hashing outperforms Robin Hood hashing in workflows with high thread counts.
- Robin Hood hashing outperforms Hopscotch hashing in update-heavy workflows or workflows with high load factors.
- All in all, Robin Hood and hopscotch are two of the most effective hashing algorithms with their own merits depending on the workflow and memory limitations.

Understanding Source Code

I. Initialization

```

7 public class HopscotchHashTable<AnyType> {
8
9     private static final int DEFAULT_TABLE_SIZE = 150;
10    private static final int RANGE = 8;
11    private HashEntry<AnyType>[] array; // The array of elements
12    private int occupied; // The number of occupied cells
13    private int theSize;
14
15    public HopscotchHashTable() {
16        this(DEFAULT_TABLE_SIZE);
17    }
18
19    public HopscotchHashTable(int size) {
20        allocateArray(size);
21        doClear();
22    }
23
24    private void allocateArray(int arraySize) {
25        array = new HashEntry[arraySize];
26    }
27
28    private void doClear() {
29        occupied = 0;
30        for (int i = 0; i < array.length; i++)
31            array[i] = null;
32    }
33
34    /**
35     * Generate hash value for the input value.
36     * @param x
37     * @return
38     */
39    private int myhash(AnyType x) {
40        int hashVal = x.hashCode();
41
42        hashVal %= array.length;
43        if (hashVal < 0)
44            hashVal += array.length;
45
46        return hashVal;
47    }

```

Here I have set default table size = 150, and the range for the “hop” = 8. Also, all the initial values are set to zero.

Hashing of the inserted element is also shown.

II. Insert

```

55 public boolean insert(AnyType x) {
56
57     if (!isEmpty()) {
58         return false;
59     }
60
61     int currentPos = findPos(x);
62
63     if (currentPos == -1) {
64         return false;
65     }
66
67     if (array[currentPos] != null) {
68         x = array[currentPos].element;
69         array[currentPos].isActive = true;
70     }
71
72     String hope;
73     if (array[currentPos] != null) {
74         hope = array[currentPos].hope;
75         x = array[currentPos].element;
76     } else {
77         hope = "10000000";
78     }
79
80     array[currentPos] = new HashEntry<>(x, hope, true);
81     theSize++;
82
83     // Rehash;
84     // if (++occupied > array.length / 2)
85     // rehash();
86     // display();
87     return true;
88 }

```

In the insert function first, we check if there is the availability of position in the table using function *isEmpty*. If there is a position left, we use *findpos* function to search for the position. If no position as per standards of hopscotch hashing is available we return false. And if the position is available we generate the neighborhood range details in a string called *hope*.

Now, for insertion of an element in the range, we add value to that block, if available, or find the new hashentry for the element

III. display

```

90 private void display() {
91     System.out.println("\n-----Display after insert-----");
92     for (int i = 0; i < array.length; i++) {
93         if (array[i] != null) {
94             System.out.println("index :" + i + " array item :"
95                 + array[i].element + " hope:" + array[i].hope);
96         } else {
97             System.out.println("index :" + i + " array item :"
98                 + "--" + " hope:" + "--");
99         }
100     }
101 }

```

Displays all the elements in the table after insertion is done

IV. contains

```

103 private boolean contains(AnyType x) {
104
105     System.out.println("\n ----- contains check for element -----:" + x);
106
107     int currentPos = myhash(x);
108     int count = 0;
109     while (array[currentPos] != null && count < RANGE) {
110
111         if (array[currentPos].hope.charAt(count) == '1') {
112             if (array[currentPos + count].element.equals(x)) {
113                 return true;
114             }
115         }
116         count++;
117     }
118
119     return false;
120 }

```

Checks if the query is there in the table using a similar method as insertion. It finds the hash-value and checks the neighborhood range as per hopscotch standards. If available, returns *true* and if not available, returns *false*

V. isEmpty

```

225 private boolean isEmpty() {
226     for (int i = 0; i < array.length; i++) {
227         if (array[i] == null) {
228             return true;
229         }
230     }
231     return false;
232 }

```

Checks if the entire table is empty or not.

VI. findPos

```

127 private int findPos(AnyType x) {
128     int offset = 0;
129     int currentPos = myhash(x);
130     int startPosition = currentPos;
131     int original = startPosition;
132     boolean flag = false;
133     boolean f = false;
134     // && !array[currentPos].element.equals(x)
135     while (array[currentPos] != null) {
136         currentPos++;
137
138         if (currentPos - startPosition >= 8
139             || (((currentPos) < startPosition) && (array.length
140                 - startPosition + currentPos) >= 8)) {
141             f = true;
142             System.out.println("flag:" + f);
143         }
144         if (f
145             && ((myhash(array[startPosition].element) - currentPos) >= RANGE
146                 || (currentPos - myhash(array[startPosition].element)) >= RANGE
147                 || ((currentPos - myhash(array[startPosition].element)) < 0 && (array.length
148                     - myhash(array[startPosition].element) + currentPos) >= 8))) {
149             flag = true;
150             currentPos = nextJumpPosition(startPosition);
151             if (currentPos == -2) {
152                 return -1;
153             }
154             startPosition = currentPos;
155             offset = 0;
156         }
157         if (currentPos >= array.length) {
158             currentPos = 0;
159         }
160     }

```

Initially setting the boolean for position to false so that it can be changed as per finding for the position according to hopscotch hashing. First, we find a position in the initial range and if not available use the **nextJumpPosition** function to find a position after shifts. Here we also handle corner cases like hash value crossing the total length of the hash table, in this case, we start from index zero.


```

161     if (flag == true) {
162         System.out.println("start:"+startPosition+"currentpos"+currentPos);
163         array[currentPos] = new HashEntry<>(array[startPosition].element,
164             "00000000", true);
165         StringBuilder string = new StringBuilder(
166             array[myhash(array[startPosition].element)].hope);
167         if ((currentPos - myhash(array[startPosition].element)) < 0) {
168             string.setCharAt(array.length
169                 - myhash(array[startPosition].element) + currentPos,
170                 '1');
171         } else {
172             string.setCharAt(currentPos
173                 - myhash(array[startPosition].element), '1');
174         }
175         array[myhash(array[startPosition].element)].hope = string
176             .toString();
177         if (array[myhash(array[startPosition].element)].hope
178             .charAt(startPosition
179                 - myhash(array[startPosition].element)) == '1') {
180             string = new StringBuilder(
181                 array[myhash(array[startPosition].element)].hope);
182             string.setCharAt(
183                 (startPosition - myhash(array[startPosition].element)),
184                 '0');
185             array[myhash(array[startPosition].element)].hope = string
186                 .toString();
187         }
188         AnyType x1 = x;
189         array[startPosition] = new HashEntry<>(x1,
190             array[startPosition].hope, true);
191         StringBuilder temp = new StringBuilder(array[myhash(x1)].hope);
192         if (startPosition - myhash(x1) < 0) {
193             temp.setCharAt(array.length - myhash(x1) + startPosition, '1');
194         } else {
195             temp.setCharAt(startPosition - myhash(x1), '1');
196         }
197         array[myhash(x1)].hope = temp.toString();
198         currentPos = startPosition;
199     } else {

```

In this part after we find a suitable position for the incoming element when the difference between the start and current position is greater than the range. Here, 8.

```

199     } else {
200         if (startPosition != currentPos) {
201             System.out.println("=" + currentPos + " start pos"
202                 + startPosition);
203             array[currentPos] = new HashEntry<>(x, "00000000", true);
204             StringBuilder temp = new StringBuilder(
205                 array[startPosition].hope);
206             int p;
207             if (currentPos - startPosition < 0) {
208                 p = array.length - startPosition + currentPos;
209                 temp.setCharAt(p, '1');
210             } else {
211                 temp.setCharAt(currentPos - startPosition, '1');
212             }
213             AnyType x1 = array[startPosition].element;
214             array[startPosition] = null;
215             array[startPosition] = new HashEntry<>(x1, temp.toString(),
216                 true);
217             currentPos = startPosition;
218         }
219     }
220 }
221 return currentPos;
222 }
223 }

```

If the above conditions are not satisfied, a new hashentry is made.

VII. nextJumpPosition

```

234- private int nextJumpPosition(int startPosition) {
235     int position = startPosition + 1;
236     int c = checkHope(position);
237     while (c == -1) {
238         if (position >= array.length) {
239             position = 0;
240             c = checkHope(position);
241         }
242         if (position == startPosition) {
243             return -2;
244         }
245         c = checkHope(position++);
246     }
247     return c;
248 }

```

Checks if the next jump is possible or not. If the next jump comes to the start position itself, the function gives a '-2' flag. If the next jump is available the function returns its position.

VIII. checkHope

```

250- private int checkHope(int position) {
251     if (array[position] != null
252         && (!array[position].hope.equals("00000000") && !array[position].hope
253             .equals("00000001"))) {
254         for (int i = 0; i < array[position].hope.length(); i++) {
255             if (array[position].hope.charAt(i) == '1') {
256                 return position + i;
257             }
258         }
259     }
260     return -1;
261 }

```

Checks the entry in the string 'hope' for the block of an array for which the function is called.

IX. Input

```

284 public static void main(String[] args) {
285     HopscotchHashTable<Integer> H = new HopscotchHashTable<>();
286     // System.out.println(H.insert(0));
287     System.out.println("insert : " + H.insert(1));
288     System.out.println("insert : " + H.insert(1));
289     System.out.println("insert : " + H.insert(3));
290     System.out.println("insert : " + H.insert(2));
291     System.out.println("insert : " + H.insert(1));
292     System.out.println("insert : " + H.insert(1));
293     System.out.println("insert : " + H.insert(1));
294     System.out.println("insert : " + H.insert(1));
295     System.out.println("insert : " + H.insert(1));
296     System.out.println("insert : " + H.insert(1));
297     System.out.println("insert : " + H.insert(1));
298     System.out.println("insert : " + H.insert(4));
299     System.out.println("insert : " + H.insert(5));
300     System.out.println("insert : " + H.insert(6));
301     System.out.println("insert : " + H.insert(7));
302     System.out.println("insert : " + H.insert(8));
303     System.out.println("insert : " + H.insert(9));
304     System.out.println("insert : " + H.insert(10));
305     System.out.println("insert : " + H.insert(25));
306     System.out.println("insert : " + H.insert(26));
307     System.out.println("insert : " + H.insert(28));
308
309     H.display();
310
311
312     System.out.println("insert : " + H1.insert(73));
313     System.out.println("insert : " + H1.insert(73));
314     System.out.println("insert : " + H1.insert(76));
315     System.out.println("insert : " + H1.insert(75));
316     System.out.println("insert : " + H1.insert(74));
317     System.out.println("insert : " + H1.insert(72));
318     System.out.println("insert : " + H1.insert(73));
319     System.out.println("insert : " + H1.insert(73));
320     System.out.println("insert : " + H1.insert(148));
321     System.out.println("insert : " + H1.insert(148));
322     System.out.println("insert : " + H1.insert(148));
323     System.out.println("insert : " + H1.insert(148));
324     System.out.println("insert : " + H1.insert(73));
325     System.out.println("insert : " + H1.insert(73));
326     System.out.println("insert : " + H1.insert(73));
327
328     H1.display();
329     System.out.println("contains : " + H1.contains(73));
330     System.out.println("contains : " + H1.contains(122));
331

```

For the first image, enters the value to show the input method. For the second image, enters conflicting values to show input and reassignment of positions. Also finally check the values that are in the table and the ones that are not in the table.

X. Output

```

insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :false
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true

-----Display after insert-----

index :0 array item :-- hope:--
index :1 array item :1 hope:11111111
index :2 array item :1 hope:00000001
index :3 array item :1 hope:00000001
index :4 array item :1 hope:00000001
index :5 array item :1 hope:00000001
index :6 array item :1 hope:00000001
index :7 array item :1 hope:00000001
index :8 array item :1 hope:00000001
index :9 array item :2 hope:00000001
index :10 array item :3 hope:00000001
index :11 array item :4 hope:00000000
index :12 array item :5 hope:00000000
index :13 array item :6 hope:00000000
index :14 array item :7 hope:00000000
index :15 array item :8 hope:00000000
index :16 array item :9 hope:00000000
index :17 array item :10 hope:00000000
index :18 array item :-- hope:--
index :19 array item :-- hope:--
index :20 array item :-- hope:--
index :21 array item :-- hope:--
index :22 array item :-- hope:--
index :23 array item :-- hope:--
index :24 array item :-- hope:--
index :25 array item :25 hope:10000000
index :26 array item :26 hope:10000000

```

Initial insert done, values shown.

-----NEW INSERT-----

```
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
insert :true
```

-----Display after insert--

```
index :0 array item :148 hope:00000000
index :1 array item :148 hope:00000000
index :2 array item :-- hope:--
index :3 array item :-- hope:--
index :4 array item :-- hope:--
index :5 array item :-- hope:--
index :6 array item :-- hope:--
index :7 array item :-- hope:--
index :8 array item :-- hope:--
index :9 array item :-- hope:--
index :10 array item :-- hope:--
index :11 array item :-- hope:--
index :12 array item :-- hope:--
index :13 array item :-- hope:--
index :14 array item :-- hope:--
index :15 array item :-- hope:--
index :16 array item :-- hope:--
index :17 array item :-- hope:--
index :18 array item :-- hope:--
index :19 array item :-- hope:--
index :20 array item :-- hope:--
index :21 array item :-- hope:--
index :22 array item :-- hope:--
index :23 array item :-- hope:--
index :24 array item :-- hope:--
index :25 array item :-- hope:--
index :26 array item :-- hope:--
index :27 array item :-- hope:--
index :28 array item :-- hope:--
index :29 array item :-- hope:--
index :30 array item :-- hope:--
index :31 array item :-- hope:--
index :32 array item :-- hope:--
index :33 array item :-- hope:--
index :34 array item :-- hope:--
index :35 array item :-- hope:--
index :36 array item :-- hope:--
index :37 array item :-- hope:--
index :38 array item :-- hope:--
index :39 array item :-- hope:--
index :40 array item :-- hope:--
index :41 array item :-- hope:--
index :42 array item :-- hope:--
index :43 array item :-- hope:--
index :44 array item :-- hope:--
index :45 array item :-- hope:--
index :46 array item :-- hope:--
index :47 array item :-- hope:--
```

```

index :48 array item :-- hope:--
index :49 array item :-- hope:--
index :50 array item :-- hope:--
index :51 array item :-- hope:--
index :52 array item :-- hope:--
index :53 array item :-- hope:--
index :54 array item :-- hope:--
index :55 array item :-- hope:--
index :56 array item :-- hope:--
index :57 array item :-- hope:--
index :58 array item :-- hope:--
index :59 array item :-- hope:--
index :60 array item :-- hope:--
index :61 array item :-- hope:--
index :62 array item :-- hope:--
index :63 array item :-- hope:--
index :64 array item :-- hope:--
index :65 array item :-- hope:--
index :66 array item :-- hope:--
index :67 array item :-- hope:--
index :68 array item :-- hope:--
index :69 array item :-- hope:--
index :70 array item :-- hope:--
index :71 array item :-- hope:--
index :72 array item :72 hope:10000000
index :73 array item :73 hope:11101111
index :74 array item :73 hope:00000001
index :75 array item :73 hope:00000001
index :76 array item :76 hope:10000000
index :77 array item :73 hope:00000000
index :78 array item :73 hope:00000000
index :79 array item :73 hope:00000000
index :80 array item :73 hope:00000000
index :81 array item :74 hope:00000000
index :82 array item :75 hope:00000000
index :83 array item :-- hope:--
index :84 array item :-- hope:--
index :85 array item :-- hope:--
index :86 array item :-- hope:--
index :87 array item :-- hope:--
index :88 array item :-- hope:--
index :89 array item :-- hope:--
index :90 array item :-- hope:--
index :91 array item :-- hope:--
index :92 array item :-- hope:--
index :93 array item :-- hope:--
index :94 array item :-- hope:--

```

Final inserts

```

index :146 array item :-- hope:--
index :147 array item :-- hope:--
index :148 array item :148 hope:11110000
index :149 array item :148 hope:00000000

----- contains check for element -----:73
contains :true

----- contains check for element -----:122
contains :false

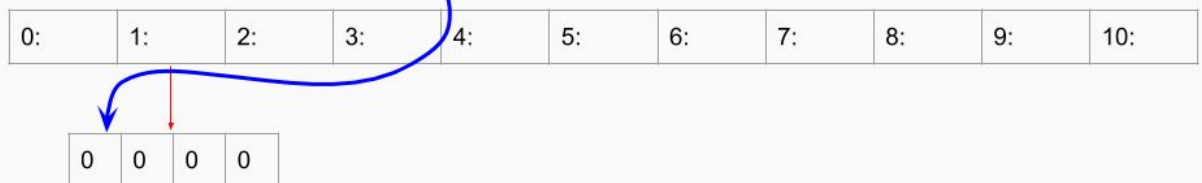
```

We can see the corner case handling, proper working of 'hope' or bit-value position storage, element finding in the table.

XI. Understanding the Concept using Graphics

Insert

Insert A \rightarrow Hash(A) \rightarrow 1 \Rightarrow A₁



Current Status A₁



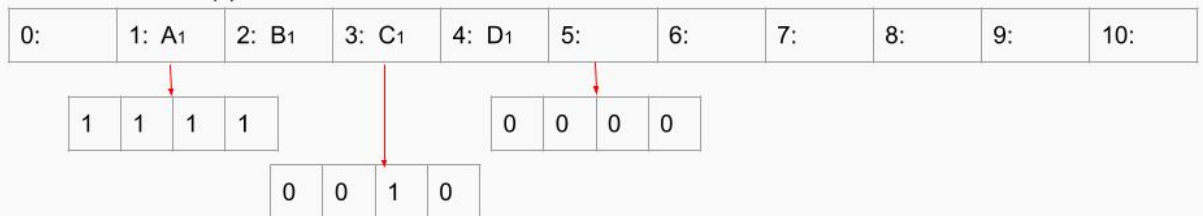
Insert B \rightarrow Hash(B) $\rightarrow 1 \Rightarrow B_1$
 Insert C \rightarrow Hash(C) $\rightarrow 1 \Rightarrow C_1$
 Insert D \rightarrow Hash(D) $\rightarrow 1 \Rightarrow D_1$



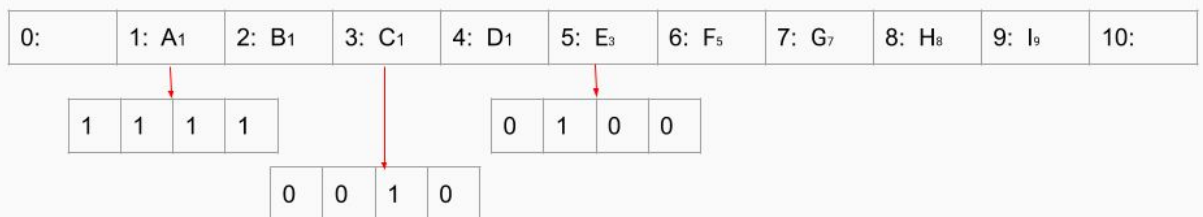
Current Status

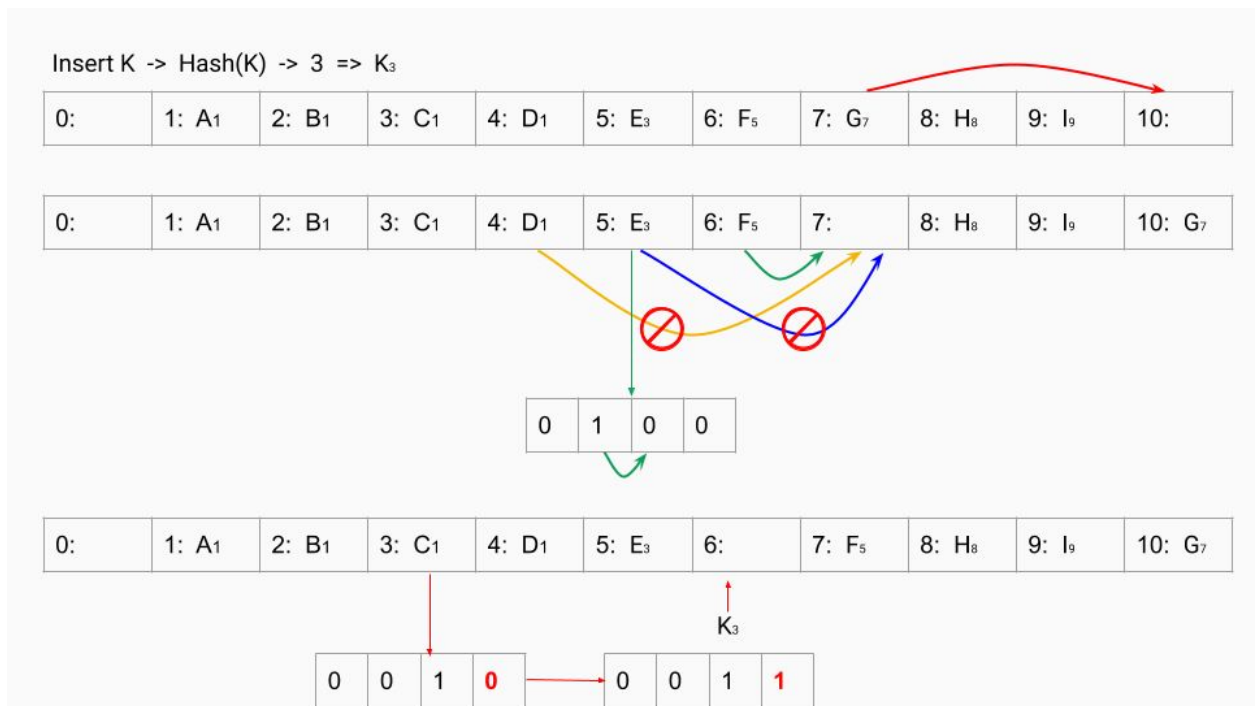


Insert E \rightarrow Hash(E) $\rightarrow 3 \Rightarrow E_3$
 Insert F \rightarrow Hash(F) $\rightarrow 5 \Rightarrow F_5$



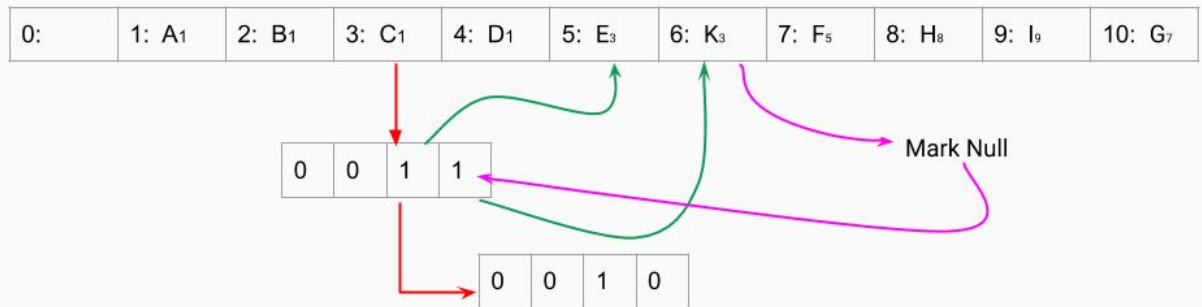
Current Status (after assume G, H, I are inserted)





Find

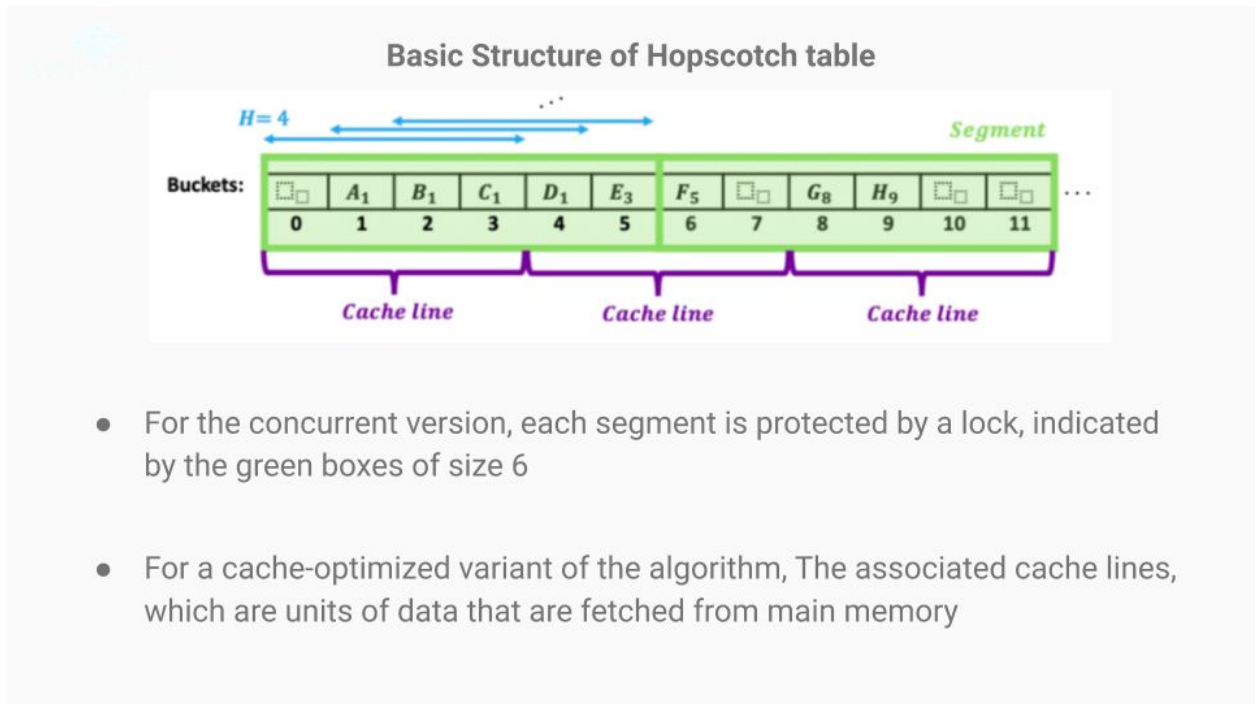
Delete K -> Find K -> Hash(K) -> 3 => K₃



What is special about this **DELETE**

- Super fast update $\sim O(1)$

Cache Support



Concurrent Hopscotch Hashing

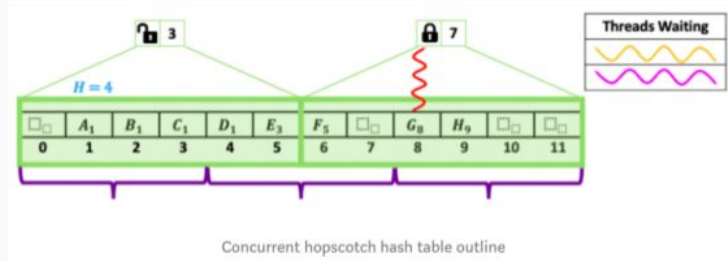


Concept

We will group buckets together into segments

Each segment will also contain its own lock and timestamp

Each time a segment is modified, that segment's timestamp is incremented by 1



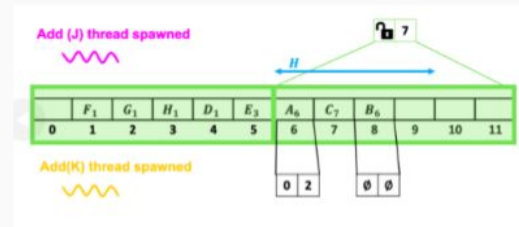
1. there cannot be any race conditions in the accessing or modification of the timestamp
2. it is possible for the contains operation to loop infinitely



Function: add()

The add operation is identical to the sequential add(k) except for two key differences:

- While traversing the hash table the relevant segment must be locked. This ensures that there are no race conditions during multiple concurrent add/remove calls to the hash table.
- Before returning, the add(k) function will increment the timestamp of the relevant segment and release the lock. The time stamp update is necessary to ensure that concurrent contains(k) re-check the updated hash table.
- worst-case complexity is $O(n)$



Bibliography

https://en.wikipedia.org/wiki/Hopscotch_hashing | Hopscotch hashing - Wikipedia

<https://programming.guide/hopscotch-hashing.html> | Hopscotch Hashing | Programming.Guide

<https://tessil.github.io/2016/08/29/hopscotch-hashing.html> | Hopscotch hashing

<http://codecapsule.com/2013/08/11/hopscotch-hashing/> | Hopscotch hashing | Code Capsule

<https://medium.com/@michelle.bao1/hopscotch-hashing-2045e7cc176b> | Overview of Hopscotch Hashing | Medium

<https://medium.com/@michelle.bao1/an-in-depth-analysis-of-hopscotch-hashings-locking-mechanisms-77b1a6f04262> | An In-Depth Analysis of Hopscotch Hashing's Locking Mechanisms | by Michelle Bao | Medium

<https://medium.com/better-programming/what-are-hash-tables-and-why-are-they-amazing-89cf52246f91> | What Are Hash Tables and Why Are They Amazing? | by Jonathan Hsu | Better Programming | Medium

<https://github.com/udand/HopscotchHashTable/blob/master/HopscotchNew/src/HopscotchHashTable.java>

Thanks!

Sridhar M
2018111021
sridhar.m@research.iiit.ac.in

